

Today in Cryptography (5830)

Hash functions

HMAC


Passwords and password-based key derivation

Where we are at

- Authenticated encryption
 - Symmetric encryption providing confidentiality and integrity
 - Security in face of active attackers
 - Uses message authentication codes as cryptographically strong error detection
 - We saw CBC-MAC, built from blockcipher
- Today: cryptographic hash functions
 - Used to build MACs, many other places
 - “Swiss army knife” of cryptography

Cryptographic hash functions

A function H that maps arbitrary bit string to fixed length string of size n



MD5: $n = 128$ bits
SHA-1: $n = 160$ bits
SHA-256: $n = 256$ bits


Many security goals asked of hash functions. Ideally, they behave as if they were a (public) random function.

Applications of hashing

- File comparison
- Digital signatures (coming up later)
- Password hashing
- For message authentication codes

Cryptographic hash functions

A function H that maps arbitrary bit string to fixed length string of size n



MD5: $n = 128$ bits

SHA-1: $n = 160$ bits

SHA-256: $n = 256$ bits

Collision resistance:

No computationally efficient adversary can find


$M \neq M'$ such that $H(M) = H(M')$

Collisions always exist

Domain (usually all strings up to some length)

SHA-1: up to length $2^{64}-1$

Range $\{0,1\}^n$




Pigeonhole principle:

size of domain larger than size of range implies there must be collisions

CBC-MAC is *not* CR

Key was secret in CBC-MAC.
But hash functions are publicly
computable.
One idea is to use a random,
public K value known to attacker.



How do we *efficiently* find
collisions?

Adversary $A(K)$:

$h \leftarrow \text{CBC-MAC}(0^n)$

$M_2 \leftarrow D_K(h) \oplus E_K(1^n)$

Return $(0^n, 1^n || M_2)$

Birthday attacks

- What is best possible security achievable by hash function with output length n bits?
- Answer: security is only achievable up to $2^{n/2}$ hash computations

The birthday problem

Choose q values Y_1, \dots, Y_q from $\{0,1\}^n$ at random. What is probability that two are the same?

Let Coll_i be event that $Y_i = Y_j$ for some $j < i$

$$\begin{aligned}\Pr[\text{Coll}] &= \Pr[\text{Coll}_1 \vee \dots \vee \text{Coll}_q] \\ &\leq \Pr[\text{Coll}_1] + \dots + \Pr[\text{Coll}_q] \\ &= \frac{0}{2^n} + \frac{1}{2^n} + \frac{2}{2^n} + \dots + \frac{q}{2^n} \\ &= \frac{q(q-1)}{2^n}\end{aligned}$$

Another proof shows that if $q \leq 2^{(n+1)/2}$

$$\Pr[\text{Coll}] \geq \frac{0.3 \cdot q(q-1)}{2^n}$$

The birthday attack

Let m be some length in domain of hash function H

Adversary A:

For $i = 1$ to q do:

$X_i \leftarrow \{0,1\}^m$

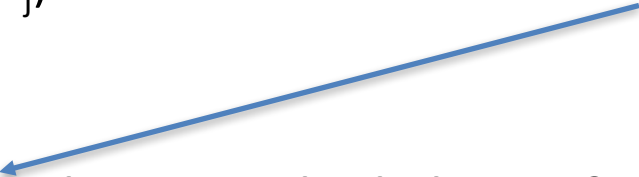
$h_i \leftarrow H(X_i)$

If exists i, j s.t. $X_i \neq X_j$ and $h_i = h_j$ then

Return (X_i, X_j)

Return fail

Same # of domain points
map to each range point



If H is *regular* then probability of success is exactly birthday probability

$$\Pr[A \text{ finds collision}] \geq \frac{0.3 \cdot q(q-1)}{2^n}$$

Birthday attack run times

MD5: $n = 128$ bits

2^{64} MD5 computations

SHA-1: $n = 160$ bits

2^{80} SHA-1 computations

SHA-256: $n = 256$ bits

2^{128} SHA-2 computations


2^{64} too small by today's standards!

Bitcoin network computes about 2^{61} SHA-256 hashes ***per second***


<https://blockchain.info/charts/hash-rate>

Two-step design for hash functions


Compression
Function



Hash Function



Domain extension called “Merkle-Damgard with strengthening”



Used in
MD-x, SHA-1,
SHA-256, ...

IV is a fixed constant. Not randomly chosen.

Building compression functions

- Can build compression functions from suitable block ciphers

$$f(z,m) = E(m,z) \oplus z$$

Called Davies-Meyer construction

- Can use AES, but security too low. Why?
- SHA-1 uses custom E with $k = d = 512$ and $n = 160$
 - Message block length of SHA-1 is 512 bits

SHA-1 compression function

Expand 512-bit message into

W_1, \dots, W_{80} strings of length 32 bit values
(Think of this as “key schedule”)

Chaining variable is 160 bits, 5 32-bit values

A, B, C, D, E

F(B,C,D) function that changes over rounds:

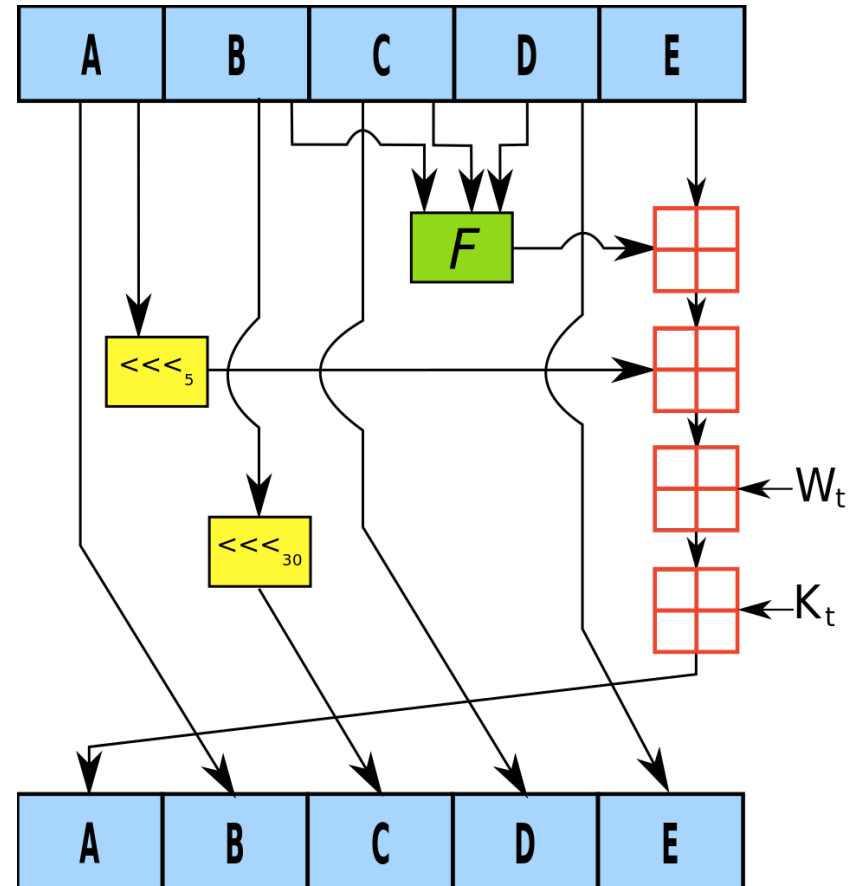
0-19: (B and C) or ((not B) and D)

20-39: $B \oplus C \oplus D$

40-59: (B and C) or (B and C) or (C and D)

60-79: $B \oplus C \oplus D$

Constants K_1, \dots, K_{80} differ across rounds



Faster attacks than birthday?

- 2004: Full break of MD5 announced by Xiaoyun Wang and co-authors
 - MD5 is easy to break now. You can download programs to do it on your laptop
- 2005: Announced faster than 2^{80} attack against SHA-1 by Wang et al.
 - Not practical to run (2^{69} estimated cost)
- 2017: CWI and Google announce first demonstrated collision

SHAttered attack

Chosen prefix P. Find two pairs of message blocks M_1, M_2 and M_1', M_2' such that for any suffix S:

$$\text{SHA-1}(P \parallel M_1 \parallel M_2 \parallel S) = \text{SHA-1}(P \parallel M_1' \parallel M_2' \parallel S)$$

Referred to as a *identical-prefix collision attack*

How? Pick P, find M_1 and M_1' that form near-collision on chaining variable. Then complete collision by finding M_2 and M_2'

They show how to extend to build colliding PDF files

SHAttered attack

SHAttered

The first concrete collision attack against SHA-1
<https://shattered.io>



Marc Stevens
Pierre Karpman



Elie Bursztein
Ange Albertini
Yarik Markov

SHAttered

The first concrete collision attack against SHA-1
<https://shattered.io>



Marc Stevens
Pierre Karpman



Elie Bursztein
Ange Albertini
Yarik Markov

Required $2^{63.1}$ SHA-1 compression function applications
100,000x faster than birthday attack (2^{80})

Fallout of attack

SVN repositories can be broken (DoS attack)

- Checking in the two SHAttered PDFs corrupts repo

Linus Torvalds misunderstands security...

(to paraphrase) GIT's ok because we can trust everyone

<https://plus.google.com/+LinusTorvalds/posts/7tp2gYWQugL>

Marc Stevens & Dan Shumow (Microsoft) developed
counter-cryptanalysis tool

Way to detect if a particular file is one half of colliding pair
Deployed at several large companies


Ongoing migration away from SHA-1 to SHA-256 / SHA-3

Applications of hashing

- File comparison
- Digital signatures (coming up later)
- Password hashing
- For message authentication codes


Building PRFs with hash functions: HMAC

Use a hash function H to build a MAC. K is a secret key




This is slight simplification, assuming $|K|$ less than block length of H
HMAC-SHA-1, HMAC-SHA-256, etc.

What's wrong with this PRF construction?



Password hashing

Password hashing. Choose random salt and store (salt,h) where:



The idea: Attacker, given (salt,h), should not be able to recover pw

Or can they?

For each guess pw':
If $H(\text{salt} || \text{pw}') = h$ then
Ret pw'

Rainbow tables speed this up in practice by way of precomputation. Large salts make rainbow tables impractical

```
rist@seclab-laptop1:~/work/teaching/642-fall-2011/slides$ openssl speed sha1
Doing sha1 for 3s on 16 size blocks: 4109047 sha1's in 3.00s
Doing sha1 for 3s on 64 size blocks: 3108267 sha1's in 2.99s
Doing sha1 for 3s on 256 size blocks: 1755265 sha1's in 3.00s
Doing sha1 for 3s on 1024 size blocks: 636540 sha1's in 3.00s
Doing sha1 for 3s on 8192 size blocks: 93850 sha1's in 3.00s
OpenSSL 1.0.0d 8 Feb 2011
```

```
rist@seclab-laptop1:~/work/teaching/642-fall-2011/slides$ openssl speed aes-128-cbc
Doing aes-128 cbc for 3s on 16 size blocks: 27022606 aes-128 cbc's in 3.00s
Doing aes-128 cbc for 3s on 64 size blocks: 6828856 aes-128 cbc's in 2.99s
Doing aes-128 cbc for 3s on 256 size blocks: 1653364 aes-128 cbc's in 3.00s
Doing aes-128 cbc for 3s on 1024 size blocks: 438909 aes-128 cbc's in 2.99s
Doing aes-128 cbc for 3s on 8192 size blocks: 54108 aes-128 cbc's in 3.00s
OpenSSL 1.0.0d 8 Feb 2011
```


Say $c = 4096$. Generous back of envelope* suggests that in 1 second, can test 252 passwords and so a naïve brute-force:

6 numerical digits	$10^6 =$ 1,000,000	~ 3968 seconds
6 lower case alphanumeric digits	$36^6 =$ 2,176,782,336	~ 99 days
8 alphanumeric + 10 special symbols	$72^8 =$ 722,204,136,308,736	~ 33million days

* I did the arithmetic...

Password-based Key Derivation (PBKDF)

PBKDF(pw,salt):



PKCS#5 standardizes PBKDF1 and PBKDF2, which are both hash-chain based.

Only slows down by a factor of c

scrypt, argon2: memory-hard hashing functions

Another application of PBKDFs: PW-based encryption

Enc(pw,M):

salt

$K \leftarrow \text{PBKDF}(\text{pw}, \text{salt})$

$C \leftarrow \text{AEnc}(K, M)$

Return (salt,C)

Here Enc is an
AE scheme
(e.g., CBC + HMAC)

Dec(pw,salt || C):

$K \leftarrow \text{PBKDF}(\text{pw}, \text{salt})$

$M \leftarrow \text{ADec}(K, C)$

Return M

Summary

- Hash functions
 - Used in a variety of applications
 - Core requirement collision resistance
- Birthday attacks break them in time $2^{n/2}$ for range size n bits
- Built from compression functions, which in turn can be viewed as block-cipher-based function
- Recent demonstration of SHA-1 collision

