

Today in Cryptography (5830)

Password hashing


Password-based AE

Where we are at

- Authenticated encryption
 - Symmetric encryption providing confidentiality and integrity
- Hash functions
 - Collision resistance
 - Use as PRFs, MACs (HMAC)
- Today:
 - Password-based key derivation
 - Password-based AE

Cryptographic hash functions

A function H that maps arbitrary bit string to fixed length string of size n



SHA-256: $n = 256$ bits

SHA-512: $n = 512$ bits

SHA-3: $n = 224, 256, 384, 512$


Many security goals asked of hash functions. Ideally, they behave as if they were a (public) random function.

Two specific security goals:

- Collision resistance
- Preimage resistance (one-wayness)

Password hashing

Password hashing. Choose random salt and store (salt,h) where:



The idea: Attacker, given (salt,h), should not be able to recover pw
A form of preimage resistance (one-wayness)

How to recover pw from (salt,h)?

For each guess pw':
If $H(\text{salt} || \text{pw}') = h$ then
Ret pw'

Rainbow tables speed this up in practice by way of precomputation. Large salts make rainbow tables impractical

Breaches are ubiquitous

⋮



32.6 million leaked (2012)
32.6 million recovered (plaintext!)



6.5 million leaked (2012)
5.85 million recovered in 2 weeks (SHA-1)



36 million accounts leaked (2013)
Encrypted, but with ECB mode




1 billion accounts (2013)
MD5 hashes

⋮

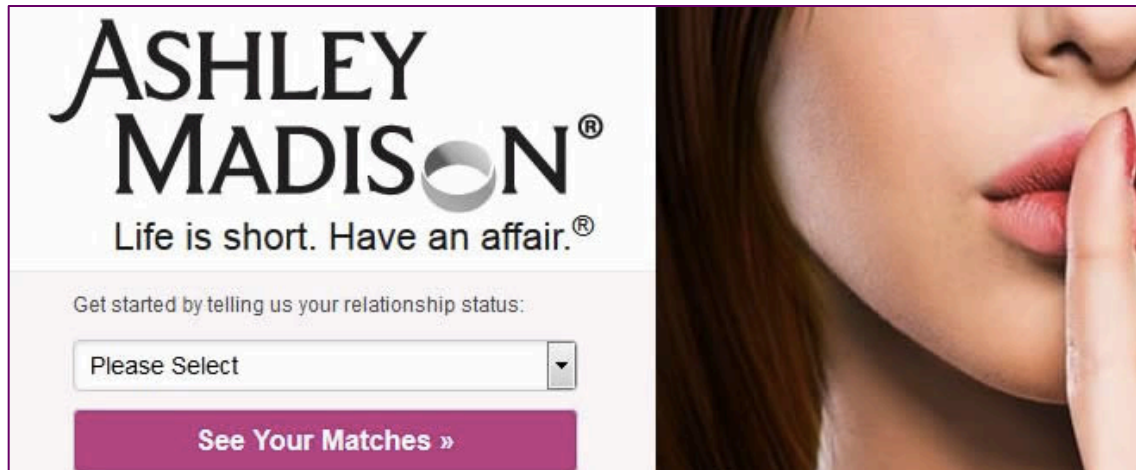
Password-based Key Derivation (PBKDF)

PBKDF(pw,salt):



PKCS#5 standardizes PBKDF1 and PBKDF2, which are both hash-chain based. (PBKDF2 uses HMAC)

Slows down cracking attacks by a factor of c



AshleyMadison hack: 36 million user hashes

Salts + Passwords hashed using bcrypt with $c = 2^{12} = 4096$
4,007 cracked directly with trivial approach

CynoSure analysis: **11 million** hashes cracked
>630,000 people used usernames as passwords
MD5 hashes left lying around accidentally

```
r1stenpart@Thomass-MacBook-Air:~/Dropbox/work/teaching/cs5830-spring2017/repo/slides$ openssl speed sha256
To get the most accurate results, try to run this
program when this computer is idle.
Doing sha256 for 3s on 16 size blocks: 4491209 sha256's in 2.98s
Doing sha256 for 3s on 64 size blocks: 2689214 sha256's in 2.98s
Doing sha256 for 3s on 256 size blocks: 1191470 sha256's in 2.99s
Doing sha256 for 3s on 1024 size blocks: 374944 sha256's in 2.98s
Doing sha256 for 3s on 8192 size blocks: 50404 sha256's in 2.99s
OpenSSL 0.9.8zg 14 July 2015
```

Say $c = 4096$. Generous back of envelope suggests that in 1 second, can test 367 passwords and so a naïve brute-force:

6 numerical digits	$10^6 =$ 1,000,000	~ 2724 seconds
6 lower case alphanumeric digits	$36^6 =$ 2,176,782,336	~ 68 days
8 alphanumeric + 10 special symbols	$72^8 =$ 722,204,136,308,736	~ 22 million days

Special-purpose hashing chips (built for Bitcoin)



14,000,000,000,000 hashes / second

Can't be used easily for password cracking, but concern about ASICs (application-specific integrated circuits) remains



Script and memory-hard hashing

- Increase time & memory needed to compute hash.
 - This makes ASIC implementations harder
- Script



Facebook password onion



```
$cur = 'password'  
$cur = md5($cur)  
$salt = randbytes(20)  
$cur = hmac_sha1($cur, $salt)  
$cur = remote_hmac_sha256($cur, $secret)  
$cur = scrypt($cur, $salt)  
$cur = hmac_sha256($cur, $salt)
```

Evolution of their password hashing over time

Limitations:

- Can't rotate secret
- Can't do cryptographic erasure for compromise clean-up

Pythia: better approach allowing secret key rotations

<http://pages.cs.wisc.edu/~ace/pythia.html>

Simple typo-tolerant password checking



Easily-corrected typos admit simple correctors:

$f_{\text{caps}}(x) = x$ with case of all letters flipped

$f_{\text{1st-case}}(x) = x$ with first letter case flipped, if it is letter

...

Define set C of corrector functions. Let $h = G_K(\text{password1})$

Relaxed checker:

If $G_K(\text{Password1}) = h$ then Return 1

For each f in C :


 If $G_K(f(\text{Password1})) = h$ then Return 1

Return 0

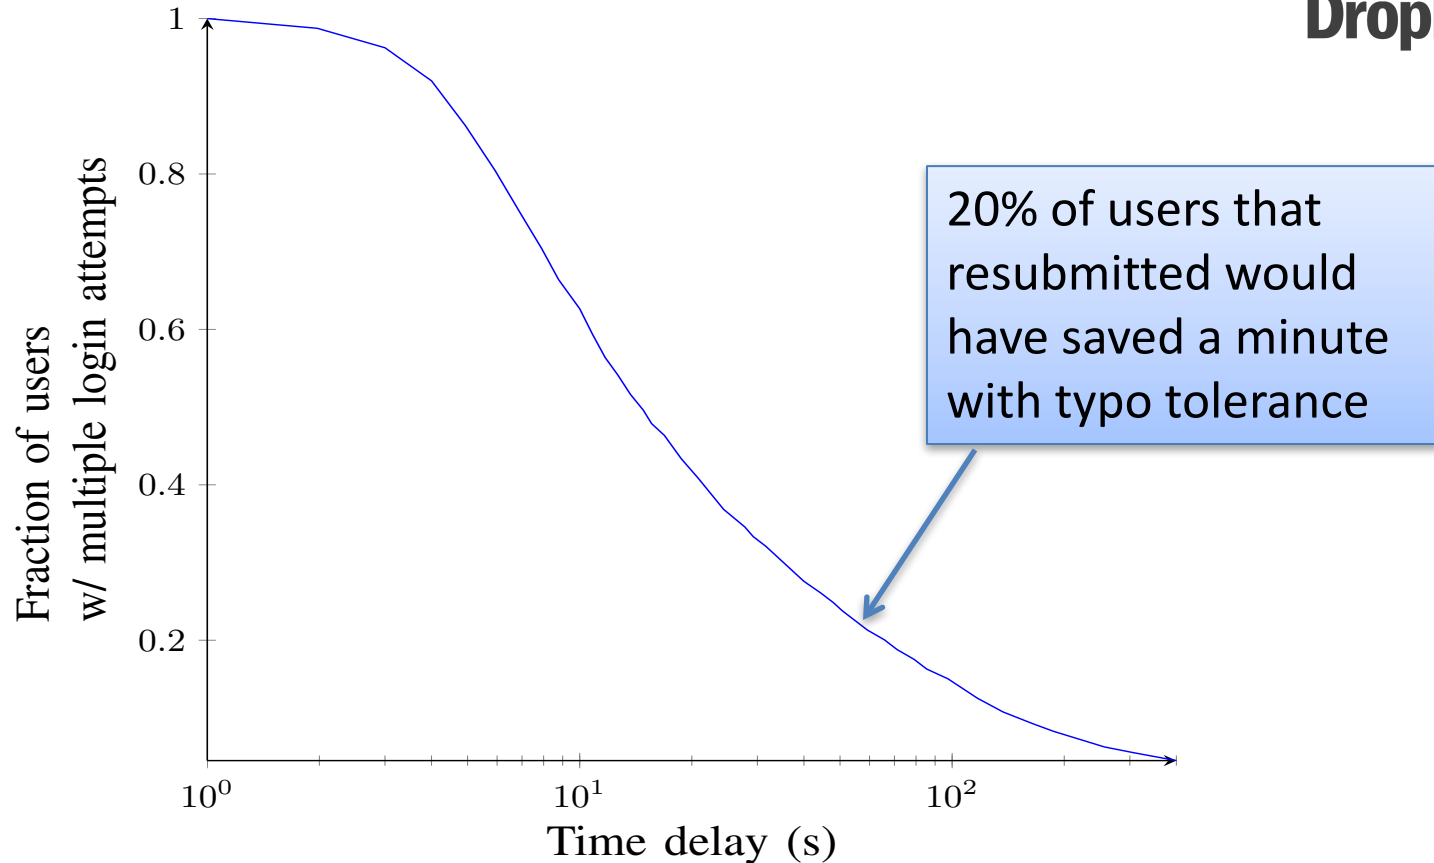
G_K is a password hashing scheme, possibly using secret K

Dropbox experiments

Utility of caps lock, first letter capitalization, removing last character



Dropbox



Typo tolerance would add several person-months of login time

3% of **users** couldn't log in during 24 hour period, but could have with one of Top 3 correctors

Empirical security analyses

Use password leaks as empirical password distributions



Typo-tolerant checker estimates distribution using RockYou



Simulate attacker following greedy strategy for typo-tolerant checker
Assume attacker knows challenge distribution exactly

q = 100	Challenge Distribution	Correctors	Exact success	Greedy strategy
	phpbb	C_{top2}	5.50%	5.50%
	phpbb	C_{top3}	5.50%	5.51%
	Myspace	C_{top2}	2.86%	2.89%
	Myspace	C_{top3}	2.86%	3.18%

Attackers that estimate challenge distribution incorrectly:
Often perform worse when trying to take advantage of tolerance
(See paper for details)

We are running new study now

- Personalized typo-tolerant system
 - Learns your specific typos over time
- We have prototype implementation for Mac, Linux

https://www.cs.cornell.edu/~rahul/projects/adaptive_typo.html

- Be great if you can participate!

Another application of PBKDFs: PW-based encryption

PWEnc(pw,M):

salt \leftarrow $\{0,1\}^{256}$

K \leftarrow PBKDF(pw,salt)

C \leftarrow Enc(K,M)

Return (salt,C)

PWDec(pw,salt || C):

K \leftarrow PBKDF(pw,salt)

M \leftarrow Dec(K,C)

Return M

Enc is a *one-time-secure*
AE scheme:

CTR-then-HMAC,
constant IV for CTR

CBC-then-HMAC,
constant IV for CBC mode

What's wrong with this?

```
import base64
from cryptography.fernet import Fernet
from cryptography.hazmat.primitives import hashes
from cryptography.hazmat.backends import default_backend

def get_key(password):

    digest = hashes.Hash(hashes.SHA256(), backend=default_backend())
    digest.update(password)
    return base64.urlsafe_b64encode(digest.finalize())

def encrypt(password, token):
    f = Fernet(get_key(key))
    return f.encrypt(bytes(token))

def decrypt(password, token):
    f = Fernet(get_key(password))
    return f.decrypt(bytes(token))
```

<http://incolumitas.com/2014/10/19/using-the-python-cryptography-module-with-custom-passwords/>

Crypto library API challenges

- Tension between *opinionated API, security and application needs*
 - Opinionated:
Low-flexibility for caller (e.g., Fernet)
 - Un-opinionated:
High-flexibility for caller (e.g., Hazmat, OpenSSL)
- Opinionated best for security, but only if it handles application needs
 - Developers will work around your API, most likely insecurely!

Cryptography.io feature request

- <https://github.com/pyca/cryptography/issues/1333>
- Need a recipe for password-based Fernet
- The game plan:
 - Paul G., myself will come up with draft spec
 - Scrypt / PBKDF + Fernet-AE
 - How to choose salt? How to serialize?
 - How to allow setting parameters (e.g., iteration count)?
 - Get feedback from Paul Kehrer, you all
 - Homework 4: individual implementations
 - Make sure they inter-operate
 - Extra credit: group project to refine spec, code

Summary

- Password hashing
 - Make hashing resource-intensive (time and/or memory)
 - Slows down brute-force cracking attacks
- Password-based authentication
 - Password onions
 - Typo-tolerance
- Password-based encryption